

MISSION

Give every developer a persistent, structural understanding of their own codebase — a real knowledge graph, not a best-effort text search — and let them bring whichever AI model they trust to reason over it. No vendor lock-in, no built-in model, no compromise on depth.

ABOUT AURA HUB

An AI-native engineering workspace (Tauri desktop app + local orchestration service) combining a codebase-grounded AI assistant, visual workflow automation, multi-step mission execution, persistent engineering memory, and rule-based governance — all bring-your-own-key, currently across 11 providers including Qwen.

Current Product

GitHub	github.com/.../aura-hub
Website	Live preview published
Documentation	Full docs/ tree, in repo
Demo	Script ready, not recorded
Architecture	1-page PDF + full docs
Production status	Feature branch, pre-main

Why Qwen

Architecture fit — Qwen's OpenAI-compatible API integrated as a 47-line adapter with zero special-casing — the cleanest proof yet of AURA's provider abstraction.

Model breadth — Alibaba publishes Qwen across multiple capability/latency tiers and long-context variants, which maps directly onto AURA's per-task model switching.

Publicly-stated strengths — Alibaba positions Qwen around coding, reasoning, and cost efficiency. We haven't independently benchmarked this yet — see the credits ask below.

Open ecosystem — Qwen's open-weight releases align with AURA's own open-source, no-lock-in philosophy.


Benefits, Both Sides

For Alibaba Cloud


Developer-tool discovery surface for Qwen; real-world engineering-platform usage data; an open feedback loop from actual integration work.

For AURA Hub


A credible, technically strong default recommendation for new users; deeper API reliability through direct feedback; a stronger enterprise story.



Real Home screen — connected environment



Qwen — pre-selected, recommended in onboarding (live)



Mission Control — real, live-created mission

Partnership Proposal

- Qwen is now the **recommended, pre-selected** provider in AURA's onboarding — first card, default selection — while full BYOAK choice is preserved for every user (no key is ever embedded; users still bring their own).
- Continue promoting Qwen through onboarding, documentation (docs/QWEN_GUIDE.md), and demo material as AURA's flagship recommendation.
- A small API credit grant to close our one honest verification gap — live authenticated inference — and produce real demo recordings.
- A technical contact to sanity-check the integration; early access to new Qwen releases; a channel for engineering feedback; consideration for any relevant developer program.

Current Status

- ✓ GitHub repository — public, Apache 2.0
- ✓ Website — published, live preview
- ✓ Documentation — full docs/ tree
- ✓ AI Runtime — 11 BYOAK providers
- ✓ Workflow Automation — 10 real templates
- ✓ Knowledge Graph — real AST analysis
- ✓ Qwen — first-class, pre-selected default
- Merged to main — pending
- Signed installers — pending

"We'd love to collaborate with Alibaba Cloud to make Qwen the default AI experience inside AURA Hub, and help grow the Qwen ecosystem together."

github.com/Aura-hub-ai-native-workspace/aura-hub